

Weighted Monte Carlo

Marius P. Furter

March 29, 2026

1 Background

The following definitions and results are included for reference.

Notation and conventions. We write $f \circ g$ for composition in *diagrammatic order*: first f , then g . We use $\llbracket \phi \rrbracket$ to denote the *Iverson bracket*: $\llbracket \phi \rrbracket = 1$ if ϕ is true and $\llbracket \phi \rrbracket = 0$ otherwise. We denote by Stoch the category of measurable spaces and probability kernels, and by sKern the category of standard Borel spaces and s-finite kernels. A *CD-category* (copy-delete category) is a symmetric monoidal category in which every object X carries a commutative comonoid structure, given by a copy map $\text{cp}_X: X \rightarrow X \otimes X$ and a delete map $\text{del}_X: X \rightarrow I$. We write $\mathcal{C}^{\text{del}}$ for the wide subcategory of arrows that commute with del . We write $\|s\|_\infty := \sup_x |s(x)|$ for the supremum norm.

Definition 1.1 (Kernels). Let (X, Σ_X) and (Y, Σ_Y) be measurable spaces. A *kernel* $f: X \rightsquigarrow Y$ is a function $f: \Sigma_Y \times X \rightarrow \mathbb{R}_{\geq 0}$ such that

- (a) for each $x \in X$, the map $B \mapsto f(B \mid x)$ is a measure,
- (b) for each $B \in \Sigma_Y$, the function $x \mapsto f(B \mid x)$ is measurable.

A kernel $f: X \rightsquigarrow Y$ is called *finite* if there is a finite uniform bound $r \in \mathbb{R}_{\geq 0}$ with $f(Y \mid x) \leq r$ for all $x \in X$. We call f *s-finite* if it can be written as a countable sum of finite kernels $f_i: X \rightsquigarrow Y$:

$$f(B \mid x) = \sum_{i=1}^{\infty} f_i(B \mid x),$$

for all $x \in X$ and $B \in \Sigma_Y$. In this case, we write $f = \sum_i f_i$ and call it a *finite decomposition*.

It is very useful to view kernels as operating on measurable functions $t: X \rightarrow \overline{\mathbb{R}}_{\geq 0}$, which we denote by $\mathfrak{M}^+(X)$.

Definition 1.2 (Kernel operator). A function $\mathbb{E}_f: \mathfrak{M}^+(Y) \rightarrow \mathfrak{M}^+(X)$ is called a *kernel operator* if it satisfies

- (a) (Linearity) For $\alpha, \beta \in \mathbb{R}_{\geq 0}$ and $s, t \in \mathfrak{M}^+(Y)$, $\mathbb{E}_f(\alpha s + \beta t) = \alpha \mathbb{E}_f(s) + \beta \mathbb{E}_f(t)$.

- (b) (Monotonicity) If $s, t \in \mathfrak{M}^+(X)$, and $s \leq t$ everywhere, then $\mathbb{E}_f(s) \leq \mathbb{E}_f(t)$.
- (c) (Monotone convergence) If t_n is a non-decreasing sequence in $\mathfrak{M}^+(Y)$, then $\lim_n \mathbb{E}_f(t_n) = \mathbb{E}_f(\lim t_n)$.

Moreover, we call $\mathbb{E}_f$ *finite* if there is $r \in \mathbb{R}_{\geq 0}$ with $\mathbb{E}_f(1) < r$ where 1 denotes the constant function. We call $\mathbb{E}_f$ *s-finite* if there is a sequence of finite kernel operators $\mathbb{E}_{f_i}$ with $\mathbb{E}_f = \sum_{i=1}^{\infty} \mathbb{E}_{f_i}$.

Theorem 1.3. *There is a one-to-one correspondence between kernels $f: X \rightsquigarrow Y$ and kernel operators $\mathbb{E}_f: \mathfrak{M}^+(Y) \rightarrow \mathfrak{M}^+(X)$. A kernel f is sent to the operator*

$$\mathbb{E}_f(t)(x) := \int_Y t(y) f(dy \mid x).$$

Conversely, each operator $\mathbb{E}_f$ is sent to the kernel

$$f(B \mid x) := \mathbb{E}_f(\llbracket y \in B \rrbracket)(x).$$

Moreover, this correspondence restricts to s-finite kernels and s-finite kernel operators. We will denote $\mathbb{E}_f[t(y) \mid x] := \mathbb{E}_f(t)(x)$.

As noted by Staton 2017 and Cho and Jacobs 2019, s-finite kernels assemble into a CD-category.

Theorem 1.4. *The category $\mathbf{sKern}$ of standard Borel spaces and s-finite kernels is a CD-category.*

The following method of constructing CD-categories is described in Fritz 2020.

Definition 1.5 (Symmetric monoidal monad). A monad M is called *symmetric monoidal* if it comes equipped with a natural transformation

$$\nabla_{X,Y}: MX \otimes MY \rightarrow M(X \otimes Y)$$

that makes M a symmetric monoidal functor with unit $\eta_I: I \rightarrow MI$, and such that η and μ are monoidal transformations. Moreover, we call M *affine* if $MI \cong I$.

Proposition 1.6. *The Kleisli category $\mathbf{Kl}_M$ of a symmetric monoidal monad $M: \mathcal{C} \rightarrow \mathcal{C}$ is a symmetric monoidal category (SMC). The monoidal product of $f_i: X_i \rightarrow MY_i$ is defined by*

$$X_1 \otimes X_2 \xrightarrow{f_1 \otimes f_2} MY_1 \otimes MY_2 \xrightarrow{\nabla} M(Y_1 \otimes Y_2).$$

Moreover, the identity-on-objects functor $L: \mathcal{C} \rightarrow \mathbf{Kl}_M$ that sends $f \mapsto f \circ \eta$ is symmetric strict monoidal.

Corollary 1.7. *Let M be a symmetric monoidal monad on a CD-category $\mathcal{C}$. Then $\mathbf{Kl}_M$ is again a CD-category. Moreover, if $\mathcal{C}$ is a Markov category and $MI \cong I$, then $\mathbf{Kl}_M$ is a Markov category as well.*

2 Weighted Monte Carlo

Weighted Monte Carlo methods approximate a kernel $f: X \rightsquigarrow Y$ by drawing samples (y_i, w_i) from a weighted probability kernel $k: X \rightsquigarrow Y \otimes W$ and forming the *weighted empirical average*

$$\mathbb{E}_f[t(y) \mid x] \approx \frac{1}{N} \sum_{i=1}^N w_i t(y_i).$$

When f is a probability kernel, we may take $w_i = 1$. More generally, the Radon-Nikodym theorem lets us sample from a convenient probability kernel q with $f \ll q$ and weight the result by the Radon-Nikodym derivative $w_i := \frac{df}{dq}(y_i \mid x)$. This is called *importance sampling*. Combining this with functoriality yields *sequential importance sampling*: given a diagram of s-finite kernels, we lift each component to a weighted probability kernel and compose; the resulting weighted samples approximate the original composite.

In practice, sequential importance sampling suffers from weight degeneracy: as the number of steps grows, the weight distribution concentrates on very few samples, wasting most computation. The standard remedy is *resampling*: maintain a population of weighted samples called *particles*, and periodically prune low-weight particles while duplicate high-weight ones. This forms the basis for *sequential Monte Carlo*.

In the following, we develop a categorical framework for sequential Monte Carlo methods. We model weighted kernels as arrows in the Kleisli category of a monoid action, and construct a symmetric strict monoidal functor that integrates weights. This structure captures sequential importance sampling via functoriality. We then extend the theory to particle populations, defining mixability conditions that characterize when resampling preserves unbiased estimates. Finally, we establish single-run convergence guarantees by analyzing the stability of resampling schemes.

2.1 Weighted kernels

2.1.1 Weights monoid

Definition 2.1 (Weights monoid). Let $\mathcal{C}$ be an SMC. A *monoid object* in $\mathcal{C}$ is an object W , along with

- (a) a multiplication arrow $m: W \otimes W \rightarrow W$,
- (b) a unit arrow $e: I \rightarrow W$,

depicted as:

$$\begin{array}{c} \text{---} \cup \text{---} \text{---} \cdot \text{---} \\ \text{---} \cup \text{---} \end{array} := m \qquad \bullet \text{---} := e$$

satisfying associativity and unitality conditions:

$$\begin{array}{c} \text{---} \cup \text{---} \cup \text{---} \text{---} \cdot \text{---} \\ \text{---} \cup \text{---} \end{array} = \begin{array}{c} \text{---} \cup \text{---} \text{---} \cup \text{---} \text{---} \cdot \text{---} \\ \text{---} \cup \text{---} \end{array} \qquad \begin{array}{c} \bullet \text{---} \cup \text{---} \text{---} \cdot \text{---} \\ \bullet \text{---} \end{array} = \text{---} = \begin{array}{c} \text{---} \cup \text{---} \text{---} \cdot \text{---} \\ \bullet \text{---} \end{array}$$

Moreover, W is called *commutative* if

From now on, we will call a commutative monoid object a *weights monoid*.

Proposition 2.2. *Let (W, m, e) be a weights monoid in an SMC $\mathcal{C}$. Then the functor $(- \otimes W): \mathcal{C} \rightarrow \mathcal{C}$ is a symmetric monoidal monad (Definition 1.5) with*

$$\begin{aligned} \mu_X: (X \otimes W) \otimes W &\xrightarrow{\alpha} X \otimes (W \otimes W) \xrightarrow{X \otimes m} X \otimes W, \\ \eta_X: X &\xrightarrow{\rho^{-1}} X \otimes I \xrightarrow{X \otimes e} X \otimes W, \\ \nabla_{X,Y}: (X \otimes W) \otimes (Y \otimes W) &\xrightarrow{s} (X \otimes Y) \otimes (W \otimes W) \xrightarrow{(X \otimes Y) \otimes m} (X \otimes Y) \otimes W, \end{aligned}$$

where s is the unique coherence arrow composed of σ and α .

Proof. This is a standard result (see e.g. Brandenburg 2014, Ex. 6.3.12). ■

We can now invoke Proposition 1.6 and Corollary 1.7 to construct new CD-categories as the Kleisli category of $(- \otimes W)$.

Corollary 2.3. *Let $\mathcal{C}$ be an SMC. Then the Kleisli category $\mathcal{C}_W := \text{Kl}_{(- \otimes W)}$ is an SMC with a symmetric strict monoidal identity-on-objects functor $e: \mathcal{C} \rightarrow \mathcal{C}_W$ sending $f \mapsto \rho^{-1} \circ (f \otimes e)$. Moreover, if $\mathcal{C}$ is a CD-category, then so is $\mathcal{C}_W$.*

A *weighted arrow* $f: X \rightarrow Y$ in $\mathcal{C}_W$ is a $\mathcal{C}$ -arrow $f: X \rightarrow Y \otimes W$. We draw the string representing the weights in red:

Kleisli composition, identities and monoidal products can then be written as follows:

Moreover, any $\mathcal{C}$ -arrow $h: X \rightarrow Y$ can be thought of as a weighted arrow by assigning its output weight e . This is precisely the action of the inclusion functor $e: \mathcal{C} \rightarrow \mathcal{C}_W$.

We can think of the Kleisli composition as the following sampling procedure:

```

| let (y, wf) = f(x) in      # draw (y, wf) ~ f(- | x)
| let (z, wg) = g(y) in      # draw (z, wg) ~ g(- | y)
| (z, wf*wg)                  # multiply weights

```

Similarly, the monoidal product corresponds procedurally to

```

| let (y1, w1) = f1(x1) in    # draw (y1, w1) ~ f1(- | x1)
| let (y2, w2) = g2(x2) in    # draw (y2, w2) ~ f2(- | x2)
| ((y1, y2), w1*w2)           # multiply weights

```

Hence, the category $\mathcal{C}_W$ implements the usual way of handling weighted samples.

2.1.2 The CD-category of weighted s-finite kernels

We are interested primarily in sKern_W , whose arrows are *weighted s-finite kernels*. We will use either *raw weights* $W^r := (\mathbb{R}_{\geq 0}, \cdot, 1)$ or *log-weights* $W^l := (\mathbb{R}, +, 0)$. Both options are equivalent under the isomorphism of monoids sending $w \mapsto \log(w)$, with the convention that $\log(0) = -\infty$ and $\log(\infty) = \infty$. If we do not say otherwise, we will work with raw weights W^r . In software implementations, however, log-weights W^l are preferable since they are more numerically stable.

We now explicitly describe the structure sKern_W . Its objects are standard Borel spaces and its arrows are s-finite weighted kernels $f: X \rightsquigarrow Y \otimes W$ (Definition 1.1). Hence, for each $x \in X$, a weighted kernel returns a joint measure over outputs Y and weights W . Given $g: Y \rightsquigarrow Z \otimes W$, the composition $f \circ g$ acts on measurable functions $t \in \mathfrak{M}^+(Z \otimes W)$ by


$$\mathbb{E}_f [\mathbb{E}_g [t(z, w_f \cdot w_g) | y] | x]$$

The Kleisli identity $\text{id}_X^{\text{Kl}}: X \rightsquigarrow X \otimes W$ is the probability kernel which acts on $t \in \mathfrak{M}^+(X \otimes W)$ by $t(x, 1)$. The monoidal product of $f_1: X_1 \rightsquigarrow Y_1 \otimes W$ and $f_2: X_2 \rightsquigarrow Y_2 \otimes W$ acts on $t \in \mathfrak{M}^+(Y_1 \otimes Y_2 \otimes W)$ according to

$$\mathbb{E}_{f_1} [\mathbb{E}_{f_2} [t(y_1, y_2, w_1 \cdot w_2) | x_2] | x_1]$$

Finally, any s-finite kernel $f: X \rightsquigarrow Y$ lifts to the weighted kernel $e(f): X \rightsquigarrow Y \otimes W$ which acts on measurable functions $t \in \mathfrak{M}^+(Y \otimes W)$ by $\mathbb{E}_f [t(y, 1)]$.

A common way to construct weighted kernels $k: X \rightsquigarrow Y \otimes W$ is by weighting a kernel $q: X \rightsquigarrow Y$ by a deterministic function $w: X \times Y \rightarrow W$:



The result acts by $\mathbb{E}_q [t(y, w(x, y))]$.


2.1.3 Integrating weights

This section describes a method of constructing functors $\mathcal{C}_W \rightarrow \mathcal{C}$ based on intuitions about weighted kernels. Any weighted kernel $f: X \rightsquigarrow Y \otimes W$ can be converted to an unweighted kernel $\bar{f}: X \rightsquigarrow Y$ by averaging over the weights:

$$\bar{f}(A \mid x) := \int_W w f(A, dw \mid x).$$

This operation can be realized by composing f with the following special kernel.

Definition 2.4 (Integrator kernel). For any space Y , the *integrator kernel* $\int_Y: Y \otimes W \rightsquigarrow Y$ is defined by $\int_Y(A \mid y, w) := w \llbracket y \in A \rrbracket$. In string diagrams, we draw integrators as follows:




It acts on $t \in \mathfrak{M}^+(Y)$ by $t(y) \mapsto w \cdot t(y)$. Composing $f: X \rightsquigarrow Y \otimes W$ with the integrator kernel yields a weighted average:

$$(f \circ_Y \int_Y)(A \mid x) = \int_Y \int_W w \llbracket y \in A \rrbracket f(dy, dw \mid x) = \int_W w f(A, dw \mid x).$$

We now describe the abstract properties of integrator kernels that will allow us to transfer the concept to other categories. First, each integrator forms a *monad algebra* over $(- \otimes W)$. In string diagrams, this means the following *unit* and *action* laws hold:



Second, the integrators on different objects assemble into a *monoidal transformation*. This entails the following compatibility with the SMC structure:




We use these properties to define abstract integrators in any SMC.


Definition 2.5. Let W be a weights monoid in an SMC $\mathcal{C}$. A *supply of integrators* for W is a monoidal transformation $\int: (- \otimes W) \Rightarrow \text{Id}$ such that every component $\int_X: X \otimes W \rightarrow X$ is an $(- \otimes W)$ algebra.

Proposition 2.6. A supply of integrators induces a symmetric strict monoidal functor $\int_W: \mathcal{C}_W \rightarrow \mathcal{C}$ that sends $f \mapsto f \circ \int$. Moreover, the functor $e: \mathcal{C} \rightarrow \mathcal{C}_W$ from Corollary 2.3 is a section of $\int_W$.

Proof. We first check that $\int_W$ is a functor. It preserves identities due to the monad algebra unit law. It preserves composition by applying the action law, followed by naturality with $(g \circ f)$:



It is a section of $e: \mathcal{C} \rightarrow \mathcal{C}_W$ due to the algebra unit law (below left). We now verify that f_W is symmetric strict monoidal. Since it is identity-on-objects by definition, we only need to check that it strictly preserves the monoidal products and coherence arrows. The former follows from the monoidal transformation law (below right):



The coherence arrows in $\mathcal{C}_W$ are of the form $e(\alpha)$. Since f_W is a section, $f_W(e(\alpha)) = \alpha$ shows that they are preserved, along with any other structures (e.g. copy and delete) lifted in this way. ■

We now prove that the claimed properties indeed hold for our main example.

Proposition 2.7. *The integrator kernels $f_Y: Y \otimes W \rightsquigarrow Y$ of Definition 2.4 form a supply of integrators for $W^r := (\mathbb{R}_{\geq 0}, \cdot, 1)$ on sKern .*

Proof. The left-hand side of the monad unit law acts by $t(x) \mapsto w \cdot t(x) \mapsto 1 \cdot t(x) = t(x)$, which agrees with the identity. The monad action law follows since the left-hand side sends $t(x) \mapsto w_2 \cdot t(x) \mapsto w_1 \cdot w_2 \cdot t(x)$, which coincides with the action of the right-hand side. Turning to naturality, the left-hand side acts by $t(x) \mapsto w \cdot t(x) \mapsto \mathbb{E}_f[w \cdot t(x)]$. The right maps $t(x) \mapsto \mathbb{E}_f[t(x)] \mapsto w \cdot \mathbb{E}_f[t(x)]$, which coincides by linearity of the integral. The compatibility with the monoidal product follows since both sides act by $t(x, y) \mapsto w_1 w_2 \cdot t(x, y)$. Finally, we note that compatibility with the monoidal unit is a special case of the algebra unit law. ■

2.1.4 Weight equivalence

Definition 2.8 (Weight equivalence). Let W be a weights monoid in an SMC $\mathcal{C}$. We say that arrows $f, g \in \mathcal{C}_W(X, Y)$ are *weight equivalent* iff $f_W f = f_W g$. In this case, we write $f \simeq g$. In string diagrams, weight equivalence becomes

$$\boxed{f} \simeq \boxed{g} \Leftrightarrow \boxed{f} \boxed{f_W} = \boxed{g} \boxed{f_W}$$

Remark 2.9. Weight equivalence is compatible with sequential and parallel composition since it is defined by the symmetric strict monoidal functor f_W .

2.1.5 Computing with weighted kernels

We are now in the position to describe how weighted kernels can be used as a computational tool. Suppose we want to approximate an s-finite kernel $m: X \rightsquigarrow Y$. Assume we can construct a weighted *probability* kernel $k: X \rightarrow Y \otimes W$ with $f_W k = m$. The next section will show that this is always possible. Given any test function $t \in \mathfrak{M}^+(Y)$, we can write $\mathbb{E}_m[t] = m \circ t = (f_W k) \circ t = k \circ (f_Y \circ t)$. Hence, given independent samples $(y^{(i)}, w^{(i)}) \sim k(- | x)$, we can approximate $\mathbb{E}_m[t]$ using the usual Monte Carlo estimator for $(f_Y \circ t)$:

$$\mathbb{E}_m[t(y) | x] \approx \frac{1}{N} \sum_i (f_Y \circ t)(y^{(i)}, w^{(i)}) = \frac{1}{N} \sum_i w^{(i)} t(y^{(i)}).$$

The approximation converges almost surely by the law of large numbers, showing that we can represent m using weighted samples from k .

In particular, if k is constructed by weighting a probability kernel $q: X \rightsquigarrow Y$ by a deterministic weight function $w: X \times Y \rightarrow W$, then the above recovers classic **importance sampling**: We may draw samples from $k(- | x)$ by first sampling $y^{(i)} \sim q(- | x)$ and then computing the importance weight $w^{(i)} := w(x, y^{(i)})$. The resulting estimator for $\mathbb{E}_m[t]$ thus becomes $\frac{1}{N} \sum_i w(x, y^{(i)}) t(y^{(i)})$. More generally, the above method may be regarded as random-weight importance sampling.

Since $\int_W$ is a symmetric strict monoidal functor, this approach extends to diagrams of s-finite kernels. In particular, suppose $m: X \rightsquigarrow Y$ is given by a diagram of s-finite kernels $f_j: X_j \rightsquigarrow Y_j$ for which there are weighted probability kernels $k_j: X_j \rightsquigarrow Y_j \otimes W$ with $\int_W k_j = f_j$. By functoriality, composing the k_j according to the diagram gives a weighted probability kernel $k: X \rightsquigarrow Y \otimes W$ with $\int_W k = m$. We may thus again approximate $\mathbb{E}_m[t]$ by the weighted average $\frac{1}{N} \sum_i w^{(i)} t(y^{(i)})$, where $(y^{(i)}, w^{(i)}) \sim k(- | x)$. Such samples may be drawn sequentially from the k_j by viewing them as a diagram in Stoch: We go through the diagram in topological order, sampling $(y_j, w_j) \sim k_j(- | x_j)$ in turn, treating the monoid multiplication as deterministic kernels. Hence, this method is a random-weight version of **sequential importance sampling**.

2.1.6 Representing kernels by weighted probability kernels

In this section we show that any s-finite kernel can be lifted to a weighted probability kernel. Explicit lifts are easy to construct when the kernel in question has a density.

Example 2.10. Suppose $k: X \rightsquigarrow Y$ is given by a density $\ell_k: Y \times X \rightarrow \mathbb{R}_{\geq 0}$ with respect to a measure ν , that is $k(B | x) = \int_Y \ell_k(y | x) \nu(dy)$. Further suppose that $q: X \rightsquigarrow Y$ is a probability kernel with density ℓ_q with respect to ν with the property that $\ell_q = 0$ implies $\ell_k = 0$. Then we may weight q by $w := \ell_k / \ell_q$ to get a weighted probability kernel $k_q: X \rightsquigarrow Y \otimes \mathbb{R}_{\geq 0}$ which acts by $\mathbb{E}_{k_q}[t(y, w(y | x))]$. Note that this is a lift of k since $\int_W k_q$ acts by

$$\begin{aligned} \mathbb{E}_{\int_W k_q}[t(y)] &= \mathbb{E}_{k_q}[w t(y)] = \int_Y w(y | x) t(y) q(dy | x) \\ &= \int_Y \frac{\ell_k(y | x)}{\ell_q(y | x)} t(y) \ell_q(y | x) \nu(dy) \\ &= \int_Y t(y) \ell_k(y | x) \nu(dy) = k \circ t. \end{aligned}$$

This idea generalizes by using the Radon-Nikodym theorem for s-finite kernels (Vákár and Ong 2018).

Definition 2.11 (Absolute continuity for s-finite kernels). Let m be a measure on a space X . We call $U \in \Sigma_X$ a

- (a) *null set*, if $m(U) = 0$,

(b) $0\text{-}\infty$ set, if for all $V \in \Sigma_X$ with $V \subseteq U$, we have $m(V) \in \{0, \infty\}$.

Moreover, for another measure m' on X , we write

- (a) $m' \ll m$, if every null set of m is a null set of m' ,
- (b) $m' \ll^\infty m$, if $m' \ll m$ and every $0\text{-}\infty$ set of m is also a $0\text{-}\infty$ set of m' .

We use the same terminology for kernels $f: X \rightsquigarrow Y$, if the corresponding property holds for $f(- | x)$, for all $x \in X$.

Lemma 2.12. *For σ -finite measures, every $0\text{-}\infty$ set is a null set. Hence, $\ll$ and $\ll^\infty$ coincide.*

Proof. Let U be a $0\text{-}\infty$ set of a σ -finite measure m on X , and suppose $m(V) = \infty$ for a measurable $V \subseteq U$. Let $\{E_i\}_{i \in \mathbb{N}}$ be a measurable partition of X with $m(E_i) < \infty$. Then $\sum_i m(V \cap E_i) = m(V)$ implies $m(V \cap E_i) > 0$ for some i . But $V \cap E_i \subseteq U$ implies $m(V \cap E_i) = \infty$, since U is a $0\text{-}\infty$ set, which contradicts $m(V \cap E_i) \leq m(E_i) < \infty$. Therefore $m(V) = 0$. ■

Theorem 2.13 (Radon-Nikodym, Vákár and Ong 2018, Thm. 10). *Let $f, q: X \rightsquigarrow Y$ be s -finite kernels with $f \ll^\infty q$. Then there exists a function*

$$\frac{df}{dq}: Y \times X \rightarrow \overline{\mathbb{R}}_{\geq 0},$$

called the Radon-Nikodym derivative, such that $\frac{df}{dq}(- | x)$ is measurable for all $x \in X$, and

$$\mathbb{E}_f[t(y) | x] = \mathbb{E}_q\left[\frac{df}{dq}(y | x) t(y) | x\right]$$

When X is countable and discrete, or Y is standard Borel, $\frac{df}{dq}$ can be made jointly measurable in X and Y . Finally, if f is σ -finite, then $\frac{df}{dq}$ takes values in $\mathbb{R}_{\geq 0}$.

Proposition 2.14. *For every s -finite kernel $f: X \rightsquigarrow Y$ in sKern there exists a probability kernel $k: X \rightsquigarrow Y \otimes \overline{\mathbb{R}}_{\geq 0}$ with $\int_W k = f$. Moreover, when f is σ -finite, the weights can be chosen to lie in $\mathbb{R}_{\geq 0}$.*

Proof. Write $f = \sum_{i=1}^{\infty} f_i$ for finite kernels $f_i: X \rightsquigarrow Y$. Define a probability kernel

$$q(A | x) := \sum_{i=1}^{\infty} 2^{-i} \frac{f_i(A | x)}{f_i(Y | x)},$$

where we interpret $0/0$ as 1. Observe that $q(A | x) = 0$ implies $f(A | x) = 0$, and $q(A | x) = \infty$ implies $f(A | x) = \infty$. Hence, $f \ll^\infty q$ and Radon-Nikodym (Theorem 2.13) yields a measurable function $v: Y \times X \rightarrow \overline{\mathbb{R}}_{\geq 0}$ satisfying

$$\mathbb{E}_f[t(y) | x] = \mathbb{E}_q[v(y, x) t(y) | x].$$

Consider the probability kernel k that acts by $t(y, w) \mapsto \mathbb{E}_q[t(y, v(y, x)) | x]$. Then

$$\mathbb{E}_{f_W k}[t(y) | x] = \mathbb{E}_k[w \cdot t(y) | x] = \mathbb{E}_q[v(y, x) \cdot t(y) | x] = \mathbb{E}_f[t(y) | x]$$

shows that $\int_W k = f$. When f is σ -finite, the Radon-Nikodym derivative v takes values in $\mathbb{R}_{\geq 0}$. ■

2.2 Parallel weighted sampling

The sequential weighted sampling method described in Section 2.1.5 has the practical limitation that the weight distribution tends to degenerate as the number of sampling steps increases, with only very few samples having relatively high weights. Since low weight samples don't contribute towards estimating weighted averages, this means that most computation is wasted.

This problem can be addressed by *resampling*: Instead of drawing samples one at a time, we maintain a population of weighted samples, called *particles*. Periodically, we eliminate low-weight samples and duplicate high-weight samples in a consistent manner. To describe such algorithms, we extend the weighted sampling theory from above to arrays of particles.


2.2.1 Mix

It is convenient to introduce the notation $x_{1:n} := (x_1, \dots, x_n)$ for arrays.

Definition 2.15 (Mix). For every space X and $n \in \mathbb{N}$, we define the *mixing kernel* $\text{mix}: X^n \rightsquigarrow X$ by its action on $t \in \mathfrak{M}^+(X)$:


$$\mathbb{E}_{\text{mix}} [t(x) \mid x_{1:n}] := \frac{1}{n} \sum_{i=1}^n t(x_i).$$

We draw mixing as



Lemma 2.16. *Mixing satisfies the following properties:*

- (a) *Copying then mixing is the identity.*
- (b) *Mix is natural on probability kernels $f: X \rightsquigarrow Y$.*



Proof. For (a), the composite acts by $t(x) \mapsto \frac{1}{n} \sum_{i=1}^n t(x_i) \mapsto \frac{1}{n} \sum_{i=1}^n t(x) = t(x)$, which is the identity. For (b), the left-hand side acts by

$$t(y) \mapsto \mathbb{E}_{f^{\otimes n}} \left[\frac{1}{n} \sum_{i=1}^n t(y_i) \mid x_i \right] = \frac{1}{n} \sum_{i=1}^n \mathbb{E}_{f^{\otimes n}} [t(y_i) \mid x_i].$$

When f is a probability kernel,

$$\mathbb{E}_{f^{\otimes n}} [t(y_i) \mid x_i] = \int_{Y_1} \cdots \int_{Y_n} t(y_i) f(dy_1 \mid x_1) \cdots f(dy_n \mid x_n) = \int_{Y_i} t(y_i) f(dy_i \mid x_i).$$

Hence, the action reduces to $t(y) \mapsto \frac{1}{n} \sum_{i=1}^n \mathbb{E}_f [t(y_i) \mid x_i]$, coinciding with right-hand side. ■

2.2.2 Category of weighted particles

Definition 2.17. A *particle transformer* $f: (X, m) \rightsquigarrow (Y, n)$ is a kernel of type $f: (X \otimes W)^m \rightsquigarrow (Y \otimes W)^n$. We think of f as accepting an array of m weighted samples and returning a random array of n weighted samples. Moreover, we use the convention that $(X, 0) = (X \otimes W)^0 = I$.

Definition 2.18 (Mixability). For $m, n \geq 1$, a particle transformer $f: (X, m) \rightsquigarrow (Y, n)$ is called *mixable* if there exists a kernel $\bar{f}: X \rightsquigarrow Y$, called *the mean of f* , satisfying:

In terms of action on test functions $t \in \mathfrak{M}^+(Y)$, this is a kind of linearity condition.

$$\mathbb{E}_f \left[\frac{1}{n} \sum_{i=1}^n w_i t(y_i) \mid (x_j, v_j)_{j=1}^m \right] = \frac{1}{m} \sum_{j=1}^m v_j \mathbb{E}_{\bar{f}} [t(y) \mid x_j],$$

In case $n = 0$ or $m = 0$, we interpret taking the weighted average across no wires as id_I , reducing the mixability conditions to

Hence, $\rho: I \rightsquigarrow (Y, n)$ is always mixable with $\mathbb{E}_{\bar{\rho}} [t(y)] := \mathbb{E}_{\rho} \left[\frac{1}{n} \sum_{i=1}^n w_i t(y_i) \right]$.

The equality $\mathbb{E}_{\bar{\rho}} [t(y)] = \mathbb{E}_{\rho} \left[\frac{1}{n} \sum_{i=1}^n w_i t(y_i) \right]$ tells us that if we compute the empirical weighted average $\frac{1}{n} \sum_{i=1}^n w_i t(y_i)$ over the output of ρ , the result will be the same as $\mathbb{E}_{\bar{\rho}} [t(y)]$ in expectation. This result extends to all mixable particle transformers.

Lemma 2.19. If $f: (X, m) \rightsquigarrow (Y, n)$ is mixable, then its mean is uniquely given by evaluating the average $\frac{1}{n} \sum_{i=1}^n t(y_i)$ at the input sequence $(x_j, v_j) := (x, 1)$, for $1 \leq j \leq m$:

$$\mathbb{E}_{\bar{f}} [t(y) \mid x] = \mathbb{E}_f \left[\frac{1}{n} \sum_{i=1}^n w_i t(y_i) \mid (x, 1)_{j=1}^m \right].$$

Proof. Suppose $m, n \geq 1$. We pre-compose both sides of the defining property with $(\text{id}_X \otimes e) \circ \text{cp}_{X \otimes W}^m$:

The final equality follows since $\text{cp}_{\text{mix}}^m = \text{id}$ (Lemma 2.16) and $(\text{id}_X \otimes e) \circ f_X = \text{id}_X$ (monad algebra unit law). Hence, the left-hand side provides a way to compute $\bar{f}$. On test functions it acts according to the claimed formula. The formula also holds for the cases $m = 0$ by disregarding the inputs, and for $n = 0$ by interpreting weighted averages over no indices as the identity. ■


Particle transformers $f: (X, m) \rightsquigarrow (Y, n)$ and $g: (Y, n) \rightsquigarrow (Z, p)$ can be composed as kernels to obtain $(f \circledcirc g): (X, m) \rightsquigarrow (Z, p)$.

Definition 2.20 (Category of particles). We define the category of *weighted particles*, denoted Part , as follows:

- (a) Its objects are pairs (X, n) , where $X \in \text{Ob}(\text{sKern})$ and $n \in \mathbb{N}$.
- (b) An arrow $f: (X, m) \rightsquigarrow (Y, n)$ is a mixable particle transformer.
- (c) Composition and identities are inherited from sKern .


Proposition 2.21. *There is a functor $P: \text{Part} \rightarrow \text{sKern}$ that sends $f \mapsto \bar{f}$.*

Proof. It is immediate from the definition of mixability that $\bar{\text{id}}_{(X \otimes W)^n} = \text{id}_X$. Next, we check that $\overline{f \circledcirc g} = \bar{f} \circledcirc \bar{g}$. The case when $m, n \geq 1$ follows from the definition:




The argument continues to work using our conventions for $m = 0$ and $n = 0$. ■

Proposition 2.22. *Given a weighted probability kernel $k: X \rightsquigarrow Y \otimes W$, we can lift it to a mixable arrow $k^n: (X, n) \rightsquigarrow (X, n)$ in Part*




whose mean is $\int_W k$. This defines a functor $(-)^n: \text{sKern}_W^{\text{del}} \rightarrow \text{Part}$.

Proof. We first argue that k^n is mixable with mean $\int_W k$ by the following calculation:



The first step uses the naturality of mix , the second the monad algebra action law, and the third the naturality of integrators. Next, we prove functoriality. The assignment preserves identities due to the monoid unit law. It preserves composition by monoid associativity and the definition of weighted kernel composition:



Corollary 2.23 (Unbiasedness of particle estimates). *The following diagram commutes:*

$$\begin{array}{ccc}
 \text{sKern}_W^{\text{del}} & \xrightarrow{(-)^n} & \text{Part} \\
 \searrow f_W & & \swarrow P \\
 & \text{sKern} &
 \end{array}$$

This justifies particle algorithms in expectation: Suppose want to approximate an s-finite kernel $f: X \rightsquigarrow Y$. As discussed above, f lifts to a weighted probability kernel $k: X \rightsquigarrow Y \otimes \mathbb{R}_{\geq 0}$ with $\int_W k = f$. By Corollary 2.23, we know $P(k^n) = \int_W k = f$. By Lemma 2.19, $\mathbb{E}_{P(k^n)} [t(y) \mid x] = \mathbb{E}_{k^n} \left[\frac{1}{n} \sum_i w_i t(y_i) \mid (x, 1)_{j=1}^m \right]$. Hence, drawing a sample $(y_{1:n}, w_{1:n})$ and computing the weighted average $\frac{1}{n} \sum_i w_i t(y_i)$ coincides with $\mathbb{E}_f [t(y)]$ in expectation. Since P and $\int_W$ are functors, this transfers to any sequential decomposition of k .

2.2.3 Resampling

The following definition abstracts the key property which resampling methods must have.

Definition 2.24. A mixable particle transformer $r: (X, m) \rightsquigarrow (X, n)$ is called a *resampler* if its mean is the identity $\bar{r} = \text{id}_X$. Explicitly, r is required to preserve weighted empirical averages in expectation:

$$\mathbb{E}_r \left[\frac{1}{n} \sum_{i=1}^n w_i t(y_i) \mid (x_j, v_j)_{j=1}^m \right] = \frac{1}{m} \sum_{j=1}^m v_j t(x_j).$$

Since $P(r) = \text{id}_X$, the functoriality of P allows us to arbitrarily insert resampling steps at the level of particle transformers without affecting the mean.

We now describe how *classic resampling schemes* fit this definition. See Chopin and Papaspiliopoulos 2020, Chapter 9 or Douc et al. 2005 for more detail. Classic schemes $r: (X, m) \rightsquigarrow (X, n)$ use the input weights $(v_j)_{j=1}^m$ to randomly choose new particles indices $(a_i)_{i=1}^n$, returning $y_i := x_{a_i}$ along with a fixed deterministic weight $w_i := \bar{v}$. This means that they act by

$$\mathbb{E}_r [w_i t(y_i) \mid (x_j, v_j)_{j=1}^m] = \bar{v} \sum_{j=1}^m \mathbb{P}(a_i = j \mid v_{1:m}) t(x_j).$$

By linearity we thus obtain

$$\mathbb{E}_r \left[\frac{1}{n} \sum_{i=1}^n w_i t(y_i) \mid (x_j, v_j)_{j=1}^m \right] = \frac{\bar{v}}{n} \sum_{j=1}^m D_j t(x_j),$$

where $D_j := \sum_{i=1}^n \mathbb{P}(a_i = j \mid v_{1:m})$ is the expected number of descendants of x_j . All classic resampling schemes satisfy $D_j := nv_j / \sum_j v_j$. Thus, setting $\bar{v} := \sum_j v_j / m$ to the average input weight ensures that r is a resampler.

Example 2.25. *Multinomial resampling* samples each index a_j independently according to the weights: $\mathbb{P}(a_i = j \mid v_{1:m}) := v_j / \sum_j v_j$. Then $D_j = \sum_{i=1}^n \mathbb{P}(a_i = j \mid v_{1:m}) = nv_j / \sum_j v_j$, as required. In practice, multinomial samples with probabilities p_i are

drawn by plugging uniform samples $u_i \sim \text{Uniform}(0, 1)$ into the inverse CDF $F^{-1}(u) = \sum_{i=1}^n i \cdot \mathbb{I}[c_{i-1} \leq u < c_i]$, where $c_i = \sum_{l=1}^i p_l$ is the i -th cumulative probability. In other words, $F(u)$ returns the largest index i for which $u < \sum_{l=1}^i p_l$.

Multinomial resampling has the downside that it introduces noise which increases the variance of the estimate. The following resampling methods aim to reduce the amount of noise by making the process more deterministic.

Example 2.26. Let $\bar{v}_j := v_j / \sum_j v_j$ denote the normalized weights. *Residual resampling* first creates $\lfloor n\bar{v}_j \rfloor$ deterministic copies of x_j and generates the remaining $R = n - \sum_j \lfloor n\bar{v}_j \rfloor$ labels by randomly sampling proportional to the residual weights $r_j = n\bar{v}_j - \lfloor n\bar{v}_j \rfloor$. Note that $\sum_j r_j = R$. Hence, we obtain the correct number of expected descendants:

$$D_j = \lfloor n\bar{v}_j \rfloor + R \frac{r_j}{\sum_j r_j} = \lfloor n\bar{v}_j \rfloor + r_j = n\bar{v}_j.$$

Example 2.27. Another class of resampling methods draws indices from a multinomial distribution by feeding seeds other than $u_i \sim \text{Uniform}(0, 1)$ into the inverse CDF described in Example 2.25. *Stratified resampling* draws $u_i \sim \text{Uniform}((i-1)/n, i/n)$. This ensures that at least one sample is present in each stratum. *Systematic resampling* goes one step further and draws a single $u \sim \text{Uniform}(0, 1)$ and then deterministically sets $u_i = (i-1 + u)/n$. The resulting samples u_i also have marginal distribution $\text{Uniform}((i-1)/n, i/n)$. We conclude by showing that both these schemes have the correct expected number of descendants. Recall that $c_i = \sum_{l=1}^i p_l$ denotes the i -th cumulative probability and that $p_i := v_j / \sum_j v_j$. Hence,

$$\begin{aligned} D_j &= \sum_{i=1}^n \mathbb{P}(a_i = j) = \sum_{i=1}^n \mathbb{P}(u_i \in [c_{j-1}, c_j]) \\ &= \sum_{i=1}^n n \cdot \lambda([c_{j-1}, c_j] \cap [(i-1)/n, i/n]) = n(c_j - c_{j-1}) = nv_j / \sum_j v_j, \end{aligned}$$

where λ denotes Lebesgue measure.

2.2.4 Single-run convergence

Recall from Theorem 1.3 that any kernel $f: X \rightsquigarrow Y$ can be equivalently viewed as a monotone linear operator $\mathbb{E}_f: \mathfrak{M}^+(Y) \rightarrow \mathfrak{M}^+(X)$ satisfying monotone convergence. In particular, for any $t \in \mathfrak{M}^+(Y)$, we have $\mathbb{E}_f[t(y) | x] = \int_Y t(y) f(dy | x)$. Note that for any deterministic kernel $d: X \rightsquigarrow Y$, we have $\mathbb{E}_d[t(y) | x] = t(d(x))$.

Moreover, since kernel composition is defined by composing their kernel operators, we get the *law of iterated expectations* $\mathbb{E}_{f \circ g} = \mathbb{E}_f \circ \mathbb{E}_g$ for free. Furthermore, our definition of kernel products in terms of iterated integrals directly yields the following classic property:

Lemma 2.28 (Expectation of independent products).

$$\mathbb{E}_{f_1 \otimes f_2} [t_1(y_1) \cdot t_2(y_2) \mid x_1, x_2] = \mathbb{E}_{f_1} [t_1(y_1) \mid x_1] \cdot \mathbb{E}_{f_2} [t_2(y_2) \mid x_2].$$

Proof. Use the definition $\mathbb{E}_{f_1 \otimes f_2} [t_1(y_1) t_2(y_2) \mid x_1, x_2] = \mathbb{E}_{f_1} [\mathbb{E}_{f_2} [t_1(y_1) t_2(y_2) \mid x_2] \mid x_1]$ and apply linearity. ■

Definition 2.29. Let $f: X \rightsquigarrow Y$ be a kernel. The *conditional variance* of a test function $t \in \mathfrak{M}^+(Y)$ with respect to f is defined by

$$\text{Var}_f [t(y) \mid x] := \mathbb{E}_f [t(y)^2 \mid x] - \mathbb{E}_f [t(y) \mid x]^2.$$

Remark 2.30. By linearity $\mathbb{E}_f$, one can show that $\text{Var}_f [t(y) \mid x] = \mathbb{E}_f [(t(y) - m)^2 \mid x]$, where $m := \mathbb{E}_f [t(y) \mid x]$. In other words, the variance measures the mean squared error of $t(y)$ to $\mathbb{E}_f [t(y) \mid x]$.

As a direct consequence of the linearity of expectation, we can pull out constants from variances.

Lemma 2.31. For any $\alpha \in \overline{\mathbb{R}}_{\geq 0}$, we have $\text{Var}_f [\alpha t(y) \mid x] = \alpha^2 \text{Var}_f [t(y) \mid x]$.

Moreover, the conditional variance of a composite can be decomposed as follows:

Lemma 2.32 (Law of total variance).

$$\text{Var}_{f \circ g} [t(z) \mid x] = \mathbb{E}_f [\text{Var}_g [t(z) \mid y] \mid x] + \text{Var}_f [\mathbb{E}_g [t(z) \mid y] \mid x].$$

Proof. We apply the definitions, the law of iterated expectations, and linearity to see that

$$\begin{aligned} \mathbb{E}_{f \circ g} [t(z)^2 \mid x] &= \mathbb{E}_f [\mathbb{E}_g [t(z)^2 \mid y] \mid x] \\ &= \mathbb{E}_f [\text{Var}_g [t(z) \mid y] + \mathbb{E}_g [t(z) \mid y]^2 \mid x] \\ &= \mathbb{E}_f [\text{Var}_g [t(z) \mid y] \mid x] + \mathbb{E}_f [\mathbb{E}_g [t(z) \mid y]^2 \mid x] \end{aligned}$$

Subtracting $\mathbb{E}_{f \circ g} [t(z) \mid x]^2 = \mathbb{E}_f [\mathbb{E}_g [t(z) \mid y] \mid x]^2$ yields the desired result. ■

Finally, when the kernels in question are probability kernels, we can split sums across independent products.

Lemma 2.33. Let $f_1: X_1 \rightsquigarrow Y_1$ and $f_2: X_2 \rightsquigarrow Y_2$ be probability kernels. Then

$$\text{Var}_{f_1 \otimes f_2} [t_1(y_1) + t_2(y_2) \mid x_1, x_2] = \text{Var}_{f_1} [t_1(y_1) \mid x_1] + \text{Var}_{f_2} [t_2(y_2) \mid x_2].$$

Proof. By definition, the variance is $\mathbb{E}_{f_1 \otimes f_2} [(t_1 + t_2)^2 \mid x_1, x_2] - \mathbb{E}_{f_1 \otimes f_2} [t_1 + t_2 \mid x_1, x_2]^2$. We use linearity, Lemma 2.28, and the fact that $\mathbb{E}_{f_1} [1] = \mathbb{E}_{f_2} [1] = 1$ since f_1 and f_2 are probability kernels to decompose the left summand into the following terms:

$$\mathbb{E}_{f_1} [t_1^2 \mid x_1], \quad 2 \mathbb{E}_{f_1} [t_1 \mid x_1] \mathbb{E}_{f_2} [t_2 \mid x_2], \quad \mathbb{E}_{f_2} [t_2^2 \mid x_2].$$

Similarly, the right summand decomposes into the following terms:

$$\mathbb{E}_{f_1} [t_1 \mid x_1]^2, \quad 2 \mathbb{E}_{f_1} [t_1 \mid x_1] \mathbb{E}_{f_2} [t_2 \mid x_2], \quad \mathbb{E}_{f_2} [t_2 \mid x_2]^2.$$

Gathering terms and applying the definition proves the result. ■

Definition 2.34. Given a measurable space X , we denote the set of bounded measurable functions $s: X \rightarrow \mathbb{R}_{\geq 0}$ by $\mathfrak{M}_b^+(X)$.

Definition 2.35 (Dispersion). Let $\rho: I \rightsquigarrow (X, n)$ be a particle transformer. The *dispersion* $d(\rho) \in \mathbb{R}_{\geq 0}$ is the smallest number satisfying

$$\text{Var}_\rho \left[\frac{1}{n} \sum_k^n w_k s(x_k) \right] \leq \frac{d(\rho)}{n} \|s\|_\infty^2,$$

for any bounded measurable function $s \in \mathfrak{M}_b^+(X)$. When $d(\rho) \in \mathbb{R}_{\geq 0}$, we say that ρ has *finite dispersion*.

Finite dispersion directly implies a *weak law of large numbers* via Chebychev's inequality. If μ is a measure on X and $\phi(x)$ a predicate, we abbreviate $\mu(\phi(x)) := \mu(\{x \in X : \phi(x)\})$. Then by definition $\mu(\phi(x)) = \mathbb{E}_\mu[\llbracket \phi(x) \rrbracket]$.

Proposition 2.36 (Markov / Chebychev inequalities). Let μ be a measure and $t \in \mathfrak{M}^+(X)$. For any $\epsilon > 0$,

$$\mu(t(x) \geq \epsilon) \leq \frac{\mathbb{E}_\mu[t(x)]}{\epsilon}, \quad \mu(|t(x) - \mathbb{E}_\mu[t(x)]| \geq \epsilon) \leq \frac{\text{Var}_\mu[t(x)]}{\epsilon^2}$$

Proof. Since $t \geq 0$, we have $\epsilon \llbracket t(x) \leq \epsilon \rrbracket \leq t(x)$. Hence, by monotonicity and linearity

$$\epsilon \mu(t(x) \leq \epsilon) = \epsilon \mathbb{E}_\mu[\llbracket t(x) \leq \epsilon \rrbracket] = \mathbb{E}_\mu[\epsilon \llbracket t(x) \leq \epsilon \rrbracket] \leq \mathbb{E}_\mu[t(x)].$$

Dividing by ϵ yields Markov's inequality. Chebychev's inequality follows by observing

$$\mu(|t(x) - \mathbb{E}_\mu[t(x)]| \geq \epsilon) = \mu((t(x) - \mathbb{E}_\mu[t(x)])^2 \geq \epsilon^2)$$

since the two sets coincide, and the applying Markov's inequality. ■

Definition 2.37 (Convergence in measure). Let μ be a measure on X . We say that a sequence of measurable functions $t_n \in \mathfrak{M}^+(X)$ *converges in measure* to $t \in \mathfrak{M}^+(X)$ with respect to μ if for any $\epsilon > 0$,

$$\mu(|t_n(x) - t(x)| \geq \epsilon) \rightarrow 0$$

as $n \rightarrow \infty$. In this case we write $t_n(x) \xrightarrow{\mu} t(x)$.

Theorem 2.38 (Weak LLN). Let $\rho: I \rightsquigarrow (X, n)$ be a particle transformer with finite dispersion. Then for every bounded measurable function $s \in \mathfrak{M}_b^+(X)$ we have

$$\frac{1}{n} \sum_k^n w_k s(x_k) \xrightarrow{\rho} \mathbb{E}_{\bar{\rho}}[s(x)].$$

Proof. Fix $\epsilon > 0$ and recall from Definition 2.18 that $\mathbb{E}_\rho[s(x)] = \mathbb{E}_\rho\left[\frac{1}{n} \sum_{k=1}^n w_k s(x_k)\right]$. Hence, by Chebychev's inequality and the definition of dispersion

$$\rho\left(\left|\frac{1}{n} \sum_k w_k s(x_k) - \mathbb{E}_\rho[s(x)]\right| \geq \epsilon\right) \leq \frac{1}{\epsilon^2} \text{Var}_\rho\left[\frac{1}{n} \sum_k w_k s(x_k)\right] \leq \frac{d\|s\|_\infty^2}{n\epsilon^2}.$$

The latter converges to 0 as $n \rightarrow \infty$, proving convergence in measure. $\blacksquare$

Definition 2.39 (Mean second moment). Given a particle transformer $\rho: I \rightsquigarrow (X, n)$, its *mean second moment* is $e(\rho) := \mathbb{E}_\rho\left[\frac{1}{n} \sum_{k=1}^n w_k^2\right]$.

Definition 2.40 (Matrix stability). A particle transformer $f: (X, n) \rightsquigarrow (Y, m)$ is called *matrix-stable* if there exists a 2×2 $\mathbb{R}$ -valued matrix M such that for every particle transformer $\rho: I \rightsquigarrow (X, n)$ we have

$$\begin{bmatrix} d(\rho \circ f) \\ e(\rho \circ f) \end{bmatrix} \leq M \begin{bmatrix} d(\rho) \\ e(\rho) \end{bmatrix}.$$

Hence, a matrix-stable f preserves finite dispersion and mean second moment.

Proposition 2.41. *Matrix-stable particle transformers are closed under composition.*

Proof. This follows immediately since

$$\begin{bmatrix} d(\rho \circ f \circ g) \\ e(\rho \circ f \circ g) \end{bmatrix} \leq M_g \begin{bmatrix} d(\rho \circ f) \\ e(\rho \circ f) \end{bmatrix} \leq M_g M_f \begin{bmatrix} d(\rho) \\ e(\rho) \end{bmatrix}.$$

$\blacksquare$

Proposition 2.42. *Let $k: X \rightsquigarrow Y \otimes \overline{\mathbb{R}}_{\geq 0}$ be a weighted kernel satisfying*

- (a) $A := \sup_{x \in X} \mathbb{E}_k[w \mid x] < \infty$,
- (b) $B := \sup_{x \in X} \mathbb{E}_k[w^2 \mid x] < \infty$.

Then the particle transformer $k^n: (X, n) \rightsquigarrow (Y, n)$ is matrix-stable with

$$M := \begin{bmatrix} A^2 & B \\ 0 & B \end{bmatrix}.$$

Proof. Let $\rho: I \rightsquigarrow (X, n)$ be a particle transformer with dispersion $d(\rho)$. We first bound the dispersion $d(\rho \circ k^n)$. By the law of total variance (Lemma 2.32), we can decompose the variance $\text{Var}_{\rho \circ k^n}\left[\frac{1}{n} \sum_i w_i s(y_i)\right] = V + E$, where

$$V := \text{Var}_\rho\left[\mathbb{E}_{k^n}\left[\frac{1}{n} \sum_i w_i s(y_i) \mid (x_j, v_j)_{j=1}^n\right]\right],$$

$$E := \mathbb{E}_\rho\left[\text{Var}_{k^n}\left[\frac{1}{n} \sum_i w_i s(y_i) \mid (x_j, v_j)_{j=1}^n\right]\right].$$

Since the k^n is mixable with mean $\int_W k$, and $\mathbb{E}_{f_W k}[s(y) \mid x] = \mathbb{E}_k[\tilde{w} s(y) \mid x]$, we have

$$\mathbb{E}_{k^n}\left[\frac{1}{n} \sum_i w_i s(y_i) \mid (x_j, v_j)_{j=1}^n\right] = \frac{1}{n} \sum_j v_j \mathbb{E}_{f_W k}[s(y_j) \mid x_j] = \frac{1}{n} \sum_j v_j \mathbb{E}_k[\tilde{w}_j s(y_j) \mid x_j].$$

By the definition of dispersion and using $\mathbb{E}_k [\tilde{w}_j s(y_j) \mid x_j] \leq \mathbb{E}_k [\tilde{w}_j \mid x_j] \|s\|_\infty \leq A \|s\|_\infty$,

$$V \leq \frac{d(\rho)}{n} \|\mathbb{E}_k [\tilde{w}_j s(y_j) \mid x_j]\|_\infty^2 \leq \frac{d(\rho) A^2 \|s\|_\infty^2}{n}.$$

For term E , we note that since k^n is an independent product over $(k \circ (\text{id}_Y \otimes m))$,

$$\text{Var}_{k^n} \left[\frac{1}{n} \sum_{i=1}^n w_i s(y_i) \mid (x_j, v_j)_{j=1}^n \right] = \frac{1}{n^2} \sum_{i=1}^n \text{Var}_{k \circ (\text{id}_Y \otimes m)} [w_i s(y_i) \mid x_i, v_i]$$

Since $(\text{id}_Y \otimes m)$ is deterministic,

$$\text{Var}_{k \circ (\text{id}_Y \otimes m)} [w_i s(y_i) \mid x_i, v_i] = \text{Var}_k [v_i \tilde{w}_i s(y_i) \mid x_i] = v_i^2 \text{Var}_k [\tilde{w}_i s(y_i) \mid x_i]$$

Moreover, we can bound

$$\text{Var}_k [\tilde{w}_i s(y_i) \mid x_i] \leq \mathbb{E}_k [\tilde{w}_i^2 s(y_i)^2 \mid x_i] \leq \mathbb{E}_k [\tilde{w}_i^2 \mid x_i] \|s\|_\infty^2 \leq B \|s\|_\infty^2.$$

Putting this together, and using monotonicity and linearity of expectation yields

$$E \leq \mathbb{E}_\rho \left[\frac{1}{n^2} \sum_i v_i^2 B \|s\|_\infty^2 \right] = \frac{B \|s\|_\infty^2}{n} \mathbb{E}_\rho \left[\frac{1}{n} \sum_i v_i^2 \right] = \frac{B \|s\|_\infty^2 e(\rho)}{n}.$$

Hence,

$$\text{Var}_{\rho \circ k^n} \left[\frac{1}{n} \sum_i w_i s(y_i) \right] \leq \frac{A^2 d(\rho) + B e(\rho)}{n} \|s\|_\infty^2$$

shows that $\rho \circ k^n$ has dispersion bounded by $A^2 d(\rho) + B e(\rho)$, as desired.

Next, we bound the mean second moment $e(\rho \circ k^n)$. Since k^n is an independent product of probability kernels $k \circ (\text{id}_Y \otimes m)$, we use Lemma 2.28 and $\mathbb{E}_{k \circ (\text{id}_Y \otimes m)} [1 \mid x_j, v_j] = 1$ to see

$$\mathbb{E}_{k^n} [w_i^2 \mid (x_j, v_j)_{j=1}^n] = \mathbb{E}_{k \circ (\text{id}_Y \otimes m)} [w_i^2 \mid x_i, v_i] = \mathbb{E}_k [v_i^2 \tilde{w}_i^2 \mid x_i] = v_i^2 \mathbb{E}_k [\tilde{w}_i^2 \mid x_i].$$

Using this, linearity, monotonicity, and iterated expectations, we compute

$$\begin{aligned} e(\rho \circ k^n) &= \mathbb{E}_{\rho \circ k^n} \left[\frac{1}{n} \sum_i w_i^2 \right] = \mathbb{E}_\rho \left[\mathbb{E}_{k^n} \left[\frac{1}{n} \sum_i w_i^2 \mid (x_j, v_j)_{j=1}^n \right] \right] \\ &= \mathbb{E}_\rho \left[\frac{1}{n} \sum_i v_i^2 \mathbb{E}_k [\tilde{w}_i^2 \mid x_i] \right] \leq \mathbb{E}_\rho \left[\frac{1}{n} \sum_i v_i^2 B \right] \\ &\leq B \mathbb{E}_\rho \left[\frac{1}{n} \sum_i v_i^2 \right] \leq B e(\rho). \end{aligned}$$

This proves the claim about matrix stability. ■

Proposition 2.43. *Let $r: (X, m) \rightsquigarrow (X, n)$ be any resampler with*

$$\begin{aligned} \text{Var}_r \left[\frac{1}{n} \sum_{i=1}^n w_i s(y_i) \mid (x_j, v_j)_{j=1}^m \right] &\leq \frac{A}{mn} \sum_{j=1}^m v_j^2 \|s\|_\infty^2, \\ \mathbb{E}_r \left[\frac{1}{n} \sum_{i=1}^n w_i^2 \mid (x_j, v_j)_{j=1}^m \right] &\leq \frac{B}{m} \sum_{j=1}^m v_j^2. \end{aligned}$$

Then r is $\begin{bmatrix} 1 & A \\ 0 & B \end{bmatrix}$ -stable.

Proof. Let $\rho: I \rightsquigarrow (X, n)$ be a particle transformer with dispersion $d(\rho)$. To bound the dispersion $d(\rho \circ r)$, we decompose the variance $\text{Var}_{\rho \circ r} \left[\frac{1}{n} \sum_i w_i s(y_i) \right] = V + E$ using Lemma 2.32, where $V := \text{Var}_\rho \left[\mathbb{E}_r \left[\frac{1}{n} \sum_{i=1}^n w_i s(y_i) \mid (x_j, v_j)_{j=1}^m \right] \right]$ and $E := \mathbb{E}_\rho \left[\text{Var}_r \left[\frac{1}{n} \sum_{i=1}^n w_i s(y_i) \mid (x_j, v_j)_{j=1}^m \right] \right]$. Since r is a resampler, it is mixable with the identity as its mean:

$$\mathbb{E}_r \left[\frac{1}{n} \sum_{i=1}^n w_i s(y_i) \mid (x_j, v_j)_{j=1}^m \right] = \frac{1}{m} \sum_{j=1}^m v_j s(x_j).$$

Hence, by the definition of dispersion, $V \leq \frac{d(\rho)}{n} \|s\|_\infty^2$. Moreover, by using the assumption and monotonicity

$$E \leq \mathbb{E}_\rho \left[\frac{\sum_{j=1}^m v_j^2}{mn} \|s\|_\infty^2 \right] = \frac{A e(\rho)}{n} \|s\|_\infty^2.$$

Therefore, the total variance can be bounded $V + E \leq \frac{d(\rho) + A e(\rho)}{n} \|s\|_\infty^2$, showing that the dispersion $d(\rho \circ r) \leq d(\rho) + A e(\rho)$.

We now bound the mean second moment. By iterated expectation

$$\mathbb{E}_{\rho \circ r} \left[\frac{1}{n} \sum_{j=1}^m v_j^2 \right] = \mathbb{E}_\rho \left[\mathbb{E}_r \left[\frac{1}{n} \sum_{i=1}^n w_i^2 \mid (x_j, v_j)_{j=1}^m \right] \right] \leq \mathbb{E}_\rho \left[B \frac{1}{m} \sum_{j=1}^m v_j^2 \right] = B e(\rho).$$

This proves the claim about stability. ■

Corollary 2.44. *Multinomial resampling is $\begin{bmatrix} 1 & 1 \\ 0 & 1 \end{bmatrix}$ -stable.*

Proof. Since multinomial resampling outputs the deterministic weight $\bar{v} := \frac{1}{m} \sum_{j=1}^m v_j$,

$$\mathbb{E}_r \left[\frac{1}{n} \sum_{i=1}^n w_i^2 \mid (x_j, v_j)_{j=1}^m \right] = \frac{1}{n} \sum_{i=1}^n \bar{v}^2 = \bar{v}^2 = \left(\frac{1}{m} \sum_{j=1}^m v_j \right)^2 \leq \frac{1}{m} \sum_{j=1}^m v_j^2,$$

where the last step uses Jensen's inequality. Furthermore, since multinomial resampling draws each sample index independently, it has the form $r = \text{cp} \circ \bigotimes_{i=1}^n r_i$. Thus, pulling out the deterministic weight and using Lemma 2.33,

$$\begin{aligned} \text{Var}_r \left[\frac{1}{n} \sum_{i=1}^n w_i s(y_i) \mid (x_j, v_j)_{j=1}^m \right] &= \frac{\bar{v}^2}{n} \text{Var}_{r_i} [s(y_i) \mid (x_j, v_j)_{j=1}^m] \\ &\leq \frac{\bar{v}^2}{n} \mathbb{E}_{r_i} [s(y_i)^2 \mid (x_j, v_j)_{j=1}^m] \leq \frac{\bar{v}^2}{n} \|s\|_\infty^2 \leq \frac{\sum_{j=1}^m v_j^2}{mn} \|s\|_\infty^2, \end{aligned}$$

where the last step uses Jensen's inequality. This shows that multinomial resampling satisfies the assumptions of Proposition 2.43 with $A = B = 1$. ■

Remark 2.45. Corollary 2.44 shows that any classic resampling scheme that has lower variance than multinomial resampling is also matrix-stable. Douc et al. 2005 show that

this is the case for residual and systematic resampling. Moreover, they provide a counterexample demonstrating that systematic resampling may fail to converge for certain unlucky particle orderings. Gerber et al. 2017 provide more detailed results on resampling schemes including proofs of almost sure weak convergence under continuity assumptions.

References

- Brandenburg, Martin (2014). “Tensor categorical foundations of algebraic geometry”. In: DOI: <https://doi.org/10.48550/arXiv.1410.1716>.
- Cho, Kenta and Bart Jacobs (Mar. 2019). “Disintegration and Bayesian inversion via string diagrams”. In: *Mathematical Structures in Computer Science* 29.7, pp. 938–971. ISSN: 1469-8072. DOI: [10.1017/S0960129518000488](https://doi.org/10.1017/S0960129518000488). eprint: [arXiv:1709.00322](https://arxiv.org/abs/1709.00322).
- Chopin, Nicolas and Omiros Papaspiliopoulos (2020). *An Introduction to Sequential Monte Carlo*. Springer International Publishing. ISBN: 9783030478452. DOI: [10.1007/978-3-030-47845-2](https://doi.org/10.1007/978-3-030-47845-2).
- Douc, Randal, Olivier Cappé, and Eric Moulines (2005). “Comparison of Resampling Schemes for Particle Filtering”. In: *Proceedings of the 4th International Symposium on Image and Signal Processing and Analysis*, pp. 64–69. DOI: [10.48550/ARXIV.CS/0507025](https://doi.org/10.48550/ARXIV.CS/0507025). eprint: [arXiv:cs/0507025](https://arxiv.org/abs/cs/0507025).
- Fritz, Tobias (Aug. 2020). “A synthetic approach to Markov kernels, conditional independence and theorems on sufficient statistics”. In: *Advances in Mathematics* 370, p. 107239. ISSN: 0001-8708. DOI: <https://doi.org/10.1016/j.aim.2020.107239>. eprint: [arXiv:1908.07021](https://arxiv.org/abs/1908.07021).
- Gerber, Mathieu, Nicolas Chopin, and Nick Whiteley (2017). “Negative association, ordering and convergence of resampling methods”. In: DOI: [10.48550/ARXIV.1707.01845](https://doi.org/10.48550/ARXIV.1707.01845). eprint: [arXiv:1707.01845](https://arxiv.org/abs/1707.01845).
- Staton, Sam (2017). “Commutative Semantics for Probabilistic Programming”. In: pp. 855–879. ISSN: 1611-3349. DOI: [10.1007/978-3-662-54434-1_32](https://doi.org/10.1007/978-3-662-54434-1_32).
- Vákár, Matthijs and Luke Ong (2018). “On S-Finite Measures and Kernels”. In: DOI: [10.48550/arXiv.1810.01837](https://doi.org/10.48550/arXiv.1810.01837).